

Hyperparameter Optimisation: A Comparative Study of Five Algorithms on YAHPO Gym

Hyperparameter Optimisation for Machine Learning — Assignment 1

Shivam Singh Gahlot (XXXXXXX) and Dlyaver Seitveliev (XXXXXXX)

RWTH Aachen University, Germany

{shivam.gahlot, dlyaver.seitveliev}@rwth-aachen.de

1 Introduction

We compare Random Search (RS), Grid Search (GS), Successive Halving (SH), Hyperband (HB), and Bayesian Optimisation (BO) on two YAHPO Gym surrogate benchmarks [?]: NB301 (neural architecture search, CIFAR-10) and rbv2_xgboost (XGBoost tuning, dataset 16). HB achieves the best final accuracy on NB301; on rbv2_xgboost the multi-fidelity methods converge faster without improving final quality.

2 Methods

Random Search (RS). RS [?] samples configurations uniformly at random and evaluates each at maximum fidelity, ignoring all prior observations.

Grid Search (GS). GS discretises each continuous dimension into 5 evenly-spaced values crossed with all categorical choices; NB301's 34-dimensional categorical space is handled by sampling each dimension independently.

Successive Halving (SH). SH [?] starts with $n = \lceil \eta^{s_{\max}+1} / (s_{\max} + 1) \rceil$ configurations at $r_{\min}$, keeps the top $1/\eta$ fraction per round, and promotes survivors to $r \cdot \eta$ until the highest rung below $r_{\max}$ ($\eta=3$, $s_{\max} = \lfloor \log_{\eta}(r_{\max}/r_{\min}) \rfloor$).

Hyperband (HB). HB [?] runs $s_{\max}+1$ SH brackets with ($n_s = \lceil (B/r_{\max}) \cdot \eta^s / (s+1) \rceil$, $r_s = \max(r_{\min}, r_{\max} \cdot \eta^{-s})$), hedging against the sensitivity of any single bracket to its initial candidate count.

Bayesian Optimisation (BO). BO [?] fits a GP (Matérn-5/2 + WhiteKernel) to all past evaluations and selects the next configuration by maximising Expected Improvement over 1 000 random candidates (5 random warm-up evaluations; GP and EI implemented from scratch in NumPy).

3 Experimental Setup


NB301: fidelity = epochs ($r_{\min}=1$, $r_{\max}=98$), metric = val. accuracy. **rbv2_xgboost/16:** fidelity = training fraction, metric = accuracy. Budget $B=50 \times r_{\max}$ fidelity units for all algorithms. Hyperparameter spaces accessed via ConfigSpace without the forbidden sampling methods; 10 independent random seeds per algorithm.

4 Results

NB301. All methods reach $\approx 93\%$ within 500 budget units (Figure ??) but diverge afterwards. HB achieves the highest final accuracy (94.16 ± 0.20); SH is weakest (92.79 ± 0.17) because a single bracket exhausts its candidate pool and restarts with no memory, whereas HB always has other brackets running at different fidelity levels. BO matches HB in mean (94.08) but is less consistent (std = 0.32, Figure ??); NB301's 34 categorical hyperparameters produce a ≈ 272 -dimensional one-hot vector that degrades the GP, making results highly seed-dependent. RS and GS perform nearly identically, as expected.

Table 1: Final incumbent (mean $\pm$ std, 10 seeds). Best in **bold**.

Algorithm	NB301	rbv2_xgb
Random Search	93.73 $\pm$ 0.17	0.962 $\pm$ 0.005
Grid Search	93.71 $\pm$ 0.18	0.963 $\pm$ 0.010
Succ. Halving	92.79 $\pm$ 0.17	0.968 $\pm$ 0.007
Hyperband	94.16 $\pm$ 0.20	0.967 $\pm$ 0.007
Bayesian Opt.	94.08 $\pm$ 0.32	0.966 $\pm$ 0.013

**Figure 1:** Incumbent curves (mean $\pm$ 1 std). Top: NB301. **Figure 2:** Final incumbent per seed. HB best on NB301; Bottom: rbv2. BO widest spread.

rbv2_xgboost. Final accuracies span only 0.006 across all methods, but the incumbent curves show SH and HB reaching ≈ 0.96 after 10 budget units while RS, GS, and BO need 20–30. XGBoost rankings at small training fractions correlate well with full-data rankings, so cheap early evaluations reliably screen out poor configurations.

5 Conclusions

HB achieves the best final accuracy on NB301 and matches SH on *rbv2_xgboost*, with lower variance than BO. Its bracket portfolio hedges against a poor initial candidate count, the key advantage over a single SH bracket. BO is competitive on NB301 in mean but unreliable due to the high-dimensional one-hot space degrading the GP surrogate. SH and HB are the practical choice on *rbv2_xgboost*, where cheap evaluations are reliable proxies for full-budget performance.

Statement of Contributions

Shivam Singh Gahlot implemented SH, HB, and BO, ran all experiments, and wrote the report. Dlyaver Seitveliev implemented RS and GS and contributed to presentation slides and visualisations.